VectorBiTE Training 2019
Methods Workshop

Intro to Maximum Likelihood

www.vectorbite.org

Assumed Background

In this workshop, we expect that you are familiar with:

- axioms of probability and their consequences.
- conditional probability and Bayes theorem
- definition of a random variable (discrete and continuous)
- the idea of a probability distribution and likelihood

Pre-workshop reading and exercises were assigned to help you review and get you ready.

We'll do a VERY fast review of likelihoods and then practice building them and finding the MLEs analytically and with R.

Method of Moments

This is one of the easiest ways to get an estimate of a what parameters are consistent with your data.
Consider an *iid* sample of n observations of a random variable $\{x_1, \ldots, x_n\}$. You can calculate sample values of the moments of the RV from these, i.e.:

$$\bar{x} = \frac{1}{n}\sum_{i=1}^{n} x_i$$

$$s^2 = \frac{1}{n}\sum (x - \bar{x})^2$$

You estimate the parameters of a probability distribution by "matching" up the sample moments with the analytical values of the moments for your probability distribution.

Likelihoods

Recall that $f(Y_i)$ is the pmf (pdf), and it tells us the probability (density) of some yet to be observed datum Y_i given a probability distribution and its parameters. If we make many observations, $\mathbf{Y} = y_1, y_2, \ldots, y_n$, we are interested how probable it was that we obtained these data, jointly. We call this the "likelihood" of the data, and denote it as

$$\mathcal{L}(\theta; Y) = f_\theta(Y)$$

where $f_\theta(Y)$ is the pdf (or pmf) of the data interpreted as a function of θ.

Likelihoods

For instance, for binomial data:

$$\Pr(Y_i = k|\theta = p) = \binom{N}{k} p^k (1-p)^{N-k}.$$

If we have data $\mathbf{Y} = y_1, y_2, \ldots, y_n$ that are i.i.d. as binomial RVs, the probabilities multiply, and the likelihood is:

$$\mathcal{L}(\theta; Y) = \prod_{i=1}^{n} \binom{N}{y_i} p^{y_i} (1-p)^{N-y_i}.$$

Likelihoods vs. probability

"Likelihood is the hypothetical probability [density] that an event that has already occurred would yield a specific outcome. The concept differs from that of a probability in that a probability refers to the occurrence of future events, while a likelihood refers to past events with known outcomes." (1)

Further, the likelihood is a function of θ (the parameters), assuming fixed data.

—

1. Weisstein, Eric W. "Likelihood." From MathWorld–A Wolfram Web Resource.

`http://mathworld.wolfram.com/Likelihood.html`

Likelihoods

We are usually interested in relative likelihoods – e.g., is it more likely that the data we observed came from a distribution with parameters θ_1 or θ_2? Thus we only worry about the likelihood up to a constant. Further, it is often easier to work with the log-likelihood:

$$L(\theta; Y) = \ell(\theta; Y) = \log(\mathcal{L}(\theta; Y))$$

where $\log(\cdot)$ is the natural log.

Maximum Likelihood Estimators (MLEs)

We can find the parameters that are most likely to have generated our data – the maximum likelihood estimate (mle) of the parameters. To do this we maximize the likelihood (or equivalently minimizing the negative log-likelihood) by taking its derivative and setting it equally to zero:

$$\frac{\partial \mathcal{L}}{\partial \theta_j} = 0 \quad \text{or} \quad -\frac{\partial L}{\partial \theta_j} = 0$$

where j denotes the j^{th} parameter. We usually denote the MLE as $\hat{\theta}_j$.

Likelihoods

The likelihood DOES NOT tell you the probability that parameters have a certain value, given the data. To obtain that quantity, usually called the "posterior probability of the parameters" in Bayesian statistics, you have to use Bayes Theorem.